

Differences between Three Levels of ANSI-SPARC Architecture

External view 1

| | | | | |
|-----|-------|-------|-----|--------|
| sNo | fName | lName | age | salary |
|-----|-------|-------|-----|--------|

External view 2

| | | |
|---------|-------|----------|
| staffNo | lName | branchNo |
|---------|-------|----------|

Conceptual level

| | | | | | |
|---------|-------|-------|-----|--------|----------|
| staffNo | fName | lName | DOB | salary | branchNo |
|---------|-------|-------|-----|--------|----------|

Internal level

```
struct STAFF {  
    int staffNo;  
    int branchNo;  
    char fName [15];  
    char lName [15];  
    struct date dateOf Birth;  
    float salary;  
    struct STAFF *next;  
};  
index staffNo; index branchNo;
```

/* pointer to next Staff record */
/* define indexes for staff */

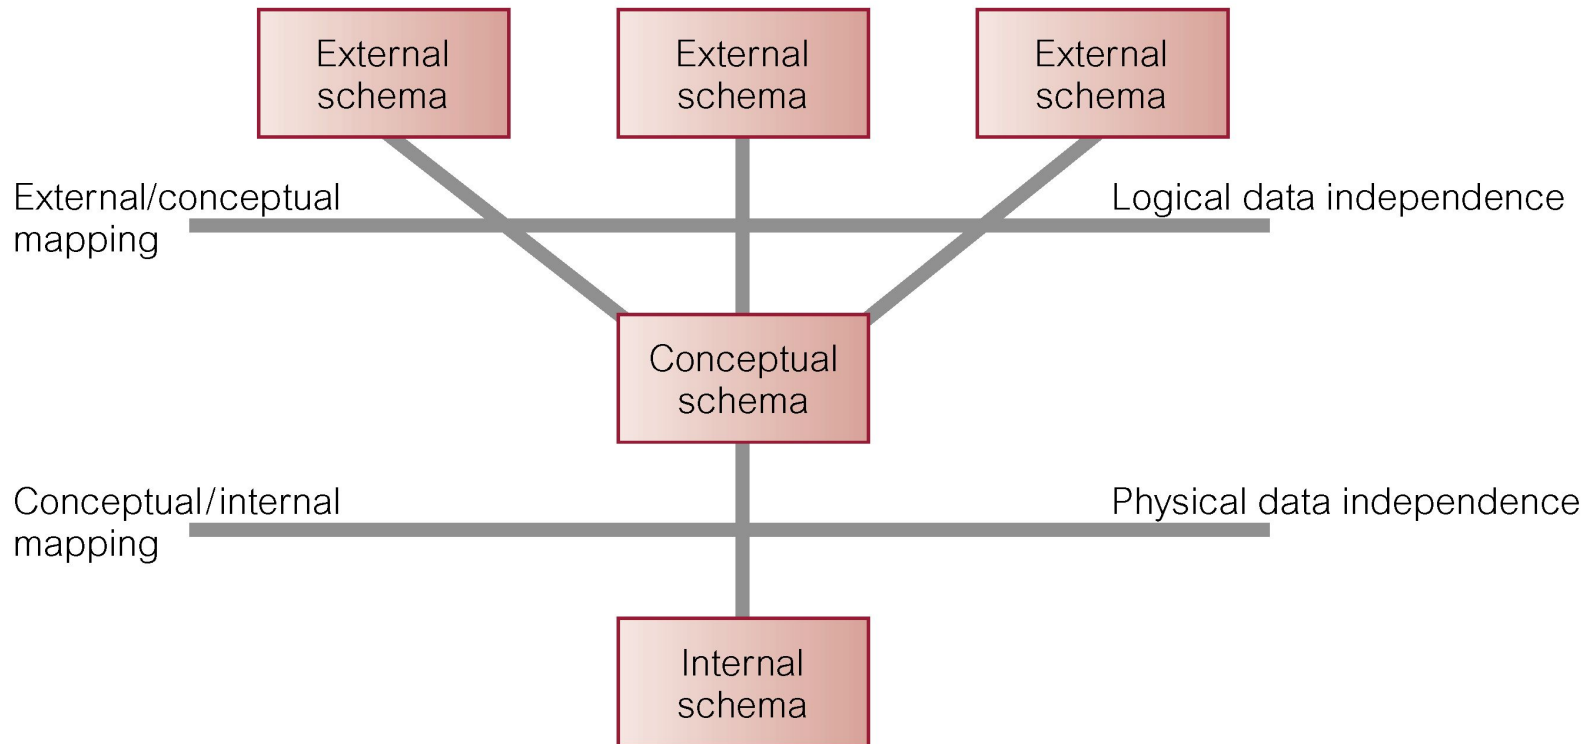
Data Independence

- Logical Data Independence
 - Refers to immunity of external schemas to changes in conceptual schema.
 - Conceptual schema changes (e.g. addition/removal of entities).
 - Should not require changes to external schema or rewrites of application programs.

Data Independence

- Physical Data Independence
 - Refers to immunity of conceptual schema to changes in the internal schema.
 - Internal schema changes (e.g. using different file organizations, storage structures/devices).
 - Should not require change to conceptual or external schemas.

Data Independence and the ANSI-SPARC Three-level Architecture



Data Model

- Data Model comprises:
 - A structural part
 - A manipulative part
 - Possibly a set of integrity rules

Data Model

- **Purpose**
 - **To represent data in an understandable way.**

Relational Systems

- Then, in 1970, E. F. Codd wrote “A Relational Model of Data for Large Shared Databanks” and introduced the relational model
- Information is stored as *tuples* or *records* in *relations* or *tables*
- Most modern DBMS are based on the relational model
- The relational model covers 3 areas:
 - Data structure
 - Data manipulation
 - Data integrity

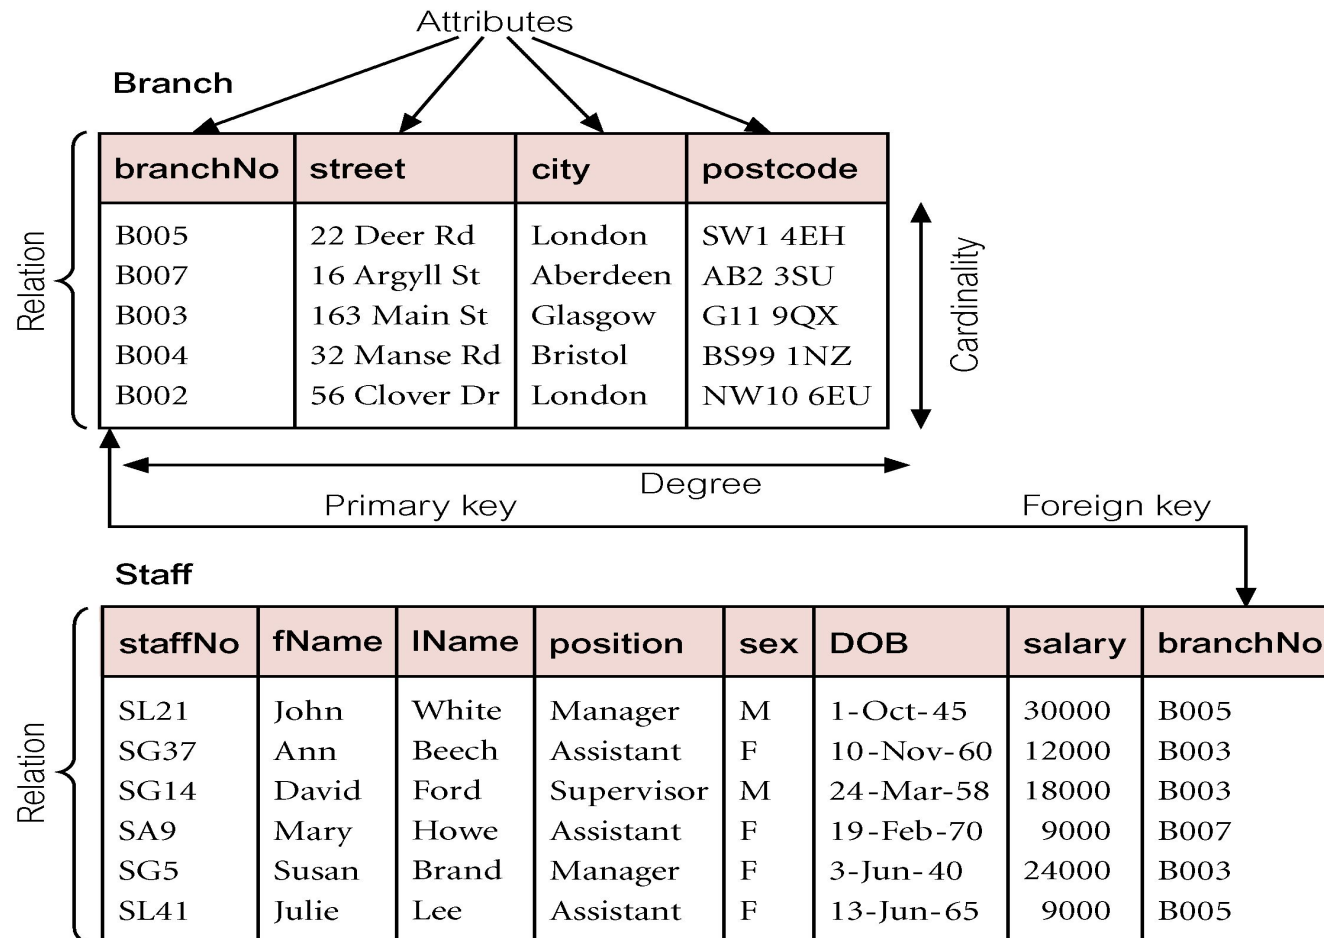
Relational Model Terminology

- A relation is a table with columns and rows.
- Attribute is a named column of a relation.
- Domain is the set of allowable values for one or more attributes.

Relational Model Terminology

- Tuple is a row of a relation.
- Degree is the number of attributes in a relation.
- Cardinality is the number of tuples in a relation.
- Relational Database is a collection of relations with distinct relation names.

Instances of Branch and Staff Relations



Examples of Attribute Domains

| Attribute | Domain Name | Meaning | Domain Definition |
|-----------|---------------|--|--|
| branchNo | BranchNumbers | The set of all possible branch numbers | character: size 4, range B001–B999 |
| street | StreetNames | The set of all street names in Britain | character: size 25 |
| city | CityNames | The set of all city names in Britain | character: size 15 |
| postcode | Postcodes | The set of all postcodes in Britain | character: size 8 |
| sex | Sex | The sex of a person | character: size 1, value M or F |
| DOB | DatesOfBirth | Possible values of staff birth dates | date, range from 1-Jan-20, format dd-mmm-yy |
| salary | Salaries | Possible values of staff salaries | monetary: 7 digits, range 6000.00–40000.00 |

Alternative Terminology for Relational Model

Table 3.1 Alternative terminology for relational model terms.

| Formal terms | Alternative 1 | Alternative 2 |
|--------------|---------------|---------------|
| Relation | Table | File |
| Tuple | Row | Record |
| Attribute | Column | Field |

Properties of Relations

- Relation name is distinct from all other relation names in relational schema.
- Each cell of relation contains exactly one atomic (single) value.
- Each attribute has a distinct name.
- Values of an attribute are all from the same domain.

Properties of Relations

- Each tuple is distinct; there are no duplicate tuples.
- Order of attributes has no significance.
- Order of tuples has no significance.

Candidate Keys

- A set of attributes in a relation is called a candidate key if, and only if,
 - Every tuple has a unique value for the set of attributes (*uniqueness*)
 - No proper subset of the set has the uniqueness property (*minimality*)

| ID | FirstLast | |
|------|-----------|-------|
| S139 | John | Smith |
| S140 | Mary | Jones |
| S141 | John | Brown |
| S142 | Jane | Smith |

Candidate key: {ID}; {First,Last} looks reasonable but we may get people with the same name

{ID, First}, {ID, Last} and {ID, First, Last} satisfy uniqueness, but are not minimal

{First} and {Last} do not give a unique identifier for each row

Choosing Candidate Keys

- Important: don't look just on the data in the table to determine what is a candidate key
- The table may contain just one tuple, so anything would do!
- Use knowledge of the real world – what is going to stay unique!

Primary Keys

- One Candidate Key is usually chosen to be used to identify tuples in a relation
- This is called the *Primary Key*
- Often a special ID attribute is used as the Primary Key

NULLs and Primary Keys

- Missing information can be represented using NULLs
- A NULL indicates a missing or unknown value
- More on this later...
- *Entity Integrity*: Primary Keys cannot contain NULL values

Foreign Keys

- *Foreign Keys* are used to link data in two relations. A set of attributes in the first (*referencing*) relation is a Foreign Key if its value always either
 - Matches a Candidate Key value in the second (*referenced*) relation, or
 - Is wholly NULL
- This is called *Referential Integrity*

Foreign Keys - Example

Department

| DID | DName |
|-----|-----------|
| 13 | Marketing |
| 14 | Accounts |
| 15 | Personnel |

{DID} is a Candidate Key for Department - Each entry has a unique value for DID

Employee

| EID | ENAME | DID |
|-----|------------|------|
| 15 | John Smith | 13 |
| 16 | Mary Brown | 14 |
| 17 | Mark Jones | 13 |
| 18 | Jane Smith | NULL |

{DID} is a Foreign Key in Employee - each Employee's DID value is either NULL, or matches an entry in the Department relation. This links each Employee to (at most) one Department

Foreign Keys - Example

Employee

| ID | Name | Manager |
|-------|------------|---------|
| E1496 | John Smith | E1499 |
| E1497 | Mary Brown | E1498 |
| E1498 | Mark Jones | E1499 |
| E1499 | Jane Smith | NULL |

{ID} is a Candidate Key for Employee, and {Manager} is a Foreign Key, which refers to the same relation - every tuple's Manager value is either NULL or matches an ID value

Referential Integrity

- When relations are updated, referential integrity can be violated
- This usually occurs when a referenced tuple is updated or deleted
- There are a number of options:
 - RESTRICT - stop the user from doing it
 - CASCADE - let the changes flow on
 - NULLIFY - make values NULL

Referential Integrity - Example

- What happens if
 - Marketing's DID is changed to 16 in Department?
 - The entry for Accounts is deleted from Department?

Department

| DID | DName |
|-----|-----------|
| 13 | Marketing |
| 14 | Accounts |
| 15 | Personnel |

Employee

| EID | ENAME | DID | |
|-----|------------|------|--|
| 15 | John Smith | 13 | |
| 16 | Mary Brown | 14 | |
| 17 | Mark Jones | 13 | |
| 18 | Jane Smith | NULL | |

RESTRICT

- RESTRICT stops any action that violates integrity
 - You cannot update or delete Marketing or Accounts
 - You *can* change Personnel as it is not referenced

Department

| DID | DName |
|-----|-----------|
| 13 | Marketing |
| 14 | Accounts |
| 15 | Personnel |

Employee

| EID | EName | DID | |
|-----|------------|------|--|
| 15 | John Smith | 13 | |
| 16 | Mary Brown | 14 | |
| 17 | Mark Jones | 13 | |
| 18 | Jane Smith | NULL | |

CASCADE

- CASCADE allows the changes made to flow through
 - If Marketing's DID is changed to 16 in Department, then the DIDs for John Smith and Mark Jones also change
 - If Accounts is deleted then so is Mary Brown

Department

| DID | DName |
|------------------|---------------------|
| 13 16 | Marketing |
| 14 | Accounts |
| 15 | Personnel |

Employee

| EID | ENAME | DID | |
|---------------|-----------------------|---------------|------------------|
| 15 | John Smith | 13 | 13 16 |
| 16 | Mary Brown | 14 | |
| 17 | Mark Jones | 13 | 13 16 |
| 18 | Jane Smith | NULL | |

NULLIFY

- NULLIFY sets problem values to NULL
 - If Marketing's DID changes then John Smith's and Mark Jones' DIDs are set to NULL
 - If Accounts is deleted, Mary Brown's DID becomes NULL

Department

| DID | DName |
|------------------|---------------------|
| 13 16 | Marketing |
| 14 | Accounts |
| 15 | Personnel |

Employee

| EID | EName | DID | |
|-----|------------|------|------|
| 15 | John Smith | 13 | NULL |
| 16 | Mary Brown | 14 | NULL |
| 17 | Mark Jones | 13 | NULL |
| 18 | Jane Smith | NULL | |

Views

- Base Relation
 - Named relation corresponding to an entity in conceptual schema, whose tuples are physically stored in database.
- View
 - Dynamic result of one or more relational operations operating on base relations to produce another relation.

Views

- A virtual relation that does not necessarily actually exist in the database but is produced upon request, at time of request.
- Contents of a view are defined as a query on one or more base relations.
- Views are dynamic, meaning that changes made to base relations that affect view attributes are immediately reflected in the view.

Purpose of Views

- Provides powerful and flexible security mechanism by hiding parts of database from certain users.
- Permits users to access data in a customized way, so that same data can be seen by different users in different ways, at same time.
- Can simplify complex operations on base relations.

Updating Views

- All updates to a base relation should be immediately reflected in all views that reference that base relation.
- If view is updated, underlying base relation should reflect change.